

```
In [ ]: import marimo as mo

In [ ]: import numpy as np
import pandas as pd
import altair as alt
alt.data_transformers.disable_max_rows()
```

DataTransformerRegistry.enable('default')

Generalized Robust Policy Optimization (GRPO)

Executive Summary

GRPO (DeepSeekMath, 2024) is a reinforcement-learning fine-tuning algorithm designed to train LLMs on verifiable tasks (math, code, reasoning) **without a value/critic model**. Key ideas at a glance:

Idea	Detail
Group advantage estimation	Sample G outputs per prompt; normalize rewards within the group
No critic needed	Baseline comes from the group mean reward — no separate value network
PPO-style clipped objective	Surrogate loss with ϵ -clipping prevents large policy updates
KL regularizer	Penalizes deviation from a reference policy, maintaining stability
Empirical gains	GSM8K: PPO 82.9% → GRPO 88.2% ; MATH: 46.8% → 51.7%

GRPO has since spawned several variants: **Dr. GRPO**, **DAPO**, **VAPO**, **Group-Robust PO**, and **Hybrid-GRPO**, each addressing different failure modes (entropy collapse, reward hacking, off-policy drift, group imbalance).

Learning Objectives

By the end of this notebook you will be able to:

1. **Explain** the motivation for GRPO and how it differs from PPO/RLHF pipelines.
2. **Derive** the group advantage estimator $A_i = (r_i - \bar{r}) / (\sigma_r + \epsilon)$.
3. **Write out** the full GRPO surrogate objective with clipping and KL penalty.
4. **Identify** the role of each hyperparameter (G , ϵ , β , μ).

5. **Compare** GRPO against PPO, TRPO, DPO, and Group-Robust PO on key axes.
6. **Interpret** the empirical benchmark results from the DeepSeekMath paper.
7. **Describe** at least three known failure modes and how GRPO variants address them.
8. **Sketch** a training loop that applies GRPO to a new reasoning task.

Background & Prerequisites

Reinforcement Learning Basics

An RL agent selects actions a_t given state s_t under policy π_θ , receiving reward r_t . The goal is to maximise expected cumulative reward $J(\theta) = \mathbb{E}_{\pi_\theta}[\sum_t r_t]$.

PPO & TRPO

TRPO (Schulman 2015) constrains the KL divergence between old and new policy to a trust region. **PPO** (Schulman 2017) approximates this via a clipped surrogate objective and a separate value network that estimates baselines. Both require a *critic* (value model) adding memory and compute overhead.

RLHF

Reinforcement Learning from Human Feedback (Ouyang 2022) trains a reward model from preference data, then runs PPO to align the LLM. The value model is a copy of the LLM — expensive at 7B+ parameters.

Robust RL

Standard RL optimises average-case reward. *Robust* RL seeks policies that perform well under worst-case or distributional perturbations — important when group-level fairness or tail-risk matters.

Mathematics Prerequisites

- Probability & expectations $\mathbb{E}[\cdot]$
- Policy gradient theorem: $\nabla_\theta J = \mathbb{E}[\nabla_\theta \log \pi_\theta(a|s) \cdot A]$
- KL divergence: $D_{\text{KL}}(p||q) = \mathbb{E}_p[\log p/q]$
- Importance sampling ratio: $\rho = \pi_\theta / \pi_{\theta_{\text{old}}}$

Core Theory

Group Advantage Estimation

For each prompt q the algorithm samples a **group** of G output sequences $\{o_1, \dots, o_G\}$ from the old policy $\pi_{\theta_{\text{old}}}$ and obtains scalar rewards $\{r_1, \dots, r_G\}$. The advantage for output i is:

$$A_i = \frac{r_i - \bar{r}}{\sigma_r + \epsilon}, \quad \bar{r} = \frac{1}{G} \sum_{j=1}^G r_j, \quad \sigma_r = \sqrt{\frac{1}{G} \sum_{j=1}^G (r_j - \bar{r})^2}$$

where $\epsilon > 0$ is a small constant for numerical stability. This *within-group normalisation* acts as a natural baseline without any learned value network.

GRPO Objective

$$J_{\text{GRPO}}(\theta) = \mathbb{E}_{q \sim \mathcal{D}, \{o_i\} \sim \pi_{\theta_{\text{old}}}(\cdot | q)} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left(\min(\rho_{i,t} A_i, \text{clip}(\rho_{i,t}, 1-\epsilon, 1+\epsilon) A_i) \right) \right]$$

where $\rho_{i,t} = \frac{\pi_{\theta}(o_{i,t} | q, o_{i,<t})}{\pi_{\theta_{\text{old}}}(o_{i,t} | q, o_{i,<t})}$ is the per-token importance-sampling ratio and β is the KL penalty weight.

Critic-Free Design

PPO must train a value head $V_{\phi}(s_t)$ to estimate $\mathbb{E}[R_t | s_t]$ and compute GAE advantages. GRPO replaces this with the *group mean* — the average reward of sibling outputs for the same prompt. This eliminates:

- A second forward/backward pass for the critic.
- The memory footprint of a value-model copy.
- Bootstrapping errors from an imperfect critic.

Mathematical Derivation

Full GRPO Objective

Expanding the per-token notation, for a single prompt q the objective is:

$$J_{\text{GRPO}}(\theta) = \frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \underbrace{\min(\rho_{i,t} A_i, \text{clip}(\rho_{i,t}, 1-\epsilon, 1+\epsilon) A_i)}_{\text{clipped surrogate}} - \underbrace{\beta \mathbb{E}_{o \sim \pi_{\theta}} \left[\log \right]}_{\text{KL per}}$$

Advantage for Outcome Supervision

With binary or scalar outcome rewards (e.g. correct/incorrect math answer):

$$A_i = \frac{r_i - \mu_r}{\sigma_r + \epsilon}, \quad r_i \in \{0, 1\} \text{ (or continuous score)}$$

PPO Advantage (GAE) vs Group-Based

	PPO (GAE)	GRPO (Group-Based)
Baseline	Learned $V_\phi(s_t)$	Group mean $\bar{r}$
Temporal credit	Discounted multi-step TD	Whole-sequence scalar
Extra model	Yes (value head)	No
Variance	Lower (if critic accurate)	Controlled by group size G
Bias	From critic errors	From finite G

The GRPO derivation follows from the policy gradient theorem applied to the group-normalised rewards, with the clipping and KL terms providing the same monotonic-improvement guarantee as PPO's trust region.

PPO vs GRPO: Key Differences

```
In [ ]: _comparison_data = pd.DataFrame({
    "Feature": [
        "Value Model", "Advantage Type", "Memory Overhead",
        "Reward Signal", "On-policy", "KL Constraint",
    ],
    "PPO": [
        "Required", "GAE (multi-step TD)", "High (2× LLM)",
        "Dense possible", "Yes", "Via penalty",
    ],
    "GRPO": [
        "Not needed", "Group-normalised", "Low (1× LLM)",
        "Outcome-level", "Yes", "Explicit  $\beta \cdot D_{KL}$ ",
    ],
})

_melted = _comparison_data.melt(id_vars="Feature", var_name="Algorithm", val

_chart = (
    alt.Chart(_melted)
    .mark_text(align="center", baseline="middle", fontSize=12)
    .encode(
        x=alt.X("Algorithm:N", axis=alt.Axis(labelFontSize=13, titleFontSize=14),
        y=alt.Y("Feature:N", axis=alt.Axis(labelFontSize=12, titleFontSize=14),
        text="Value:N",
        color=alt.Color(
            "Algorithm:N",
            scale=alt.Scale(range=["#4e79a7", "#f28e2b"]),
            legend=alt.Legend(title="Algorithm"),
        ),
    ),
)
```

```

        .properties(
            title="PPO vs GRPO Feature Comparison",
            width=520,
            height=260,
        )
        .interactive()
    )
    ppo_grpo_comparison_chart = _chart

```

```
In [ ]: ppo_grpo_comparison_chart
```

Surrogate Loss Clipping (Interactive)

```

In [ ]: _rho = np.linspace(0.5, 1.5, 200)
        _eps = 0.2
        _clipped_rho = np.clip(_rho, 1 - _eps, 1 + _eps)

        _advantage_specs = [(1.0, "A = +1 (positive)"), (-1.0, "A = -1 (negative)")]
        _records = [
            {"rho": rho_val, "value": rho_val * A, "type": "Unclipped", "advantage":
            for rho_val, clipped_val in zip(_rho, _clipped_rho)
            for A, label in _advantage_specs
        ] + [
            {"rho": rho_val, "value": clipped_val * A, "type": "Clipped", "advantage":
            for rho_val, clipped_val in zip(_rho, _clipped_rho)
            for A, label in _advantage_specs
        ]

        _clip_df = pd.DataFrame(_records)

        _region = pd.DataFrame({"x1": [0.8], "x2": [1.2]})
        _region_band = (
            alt.Chart(_region)
            .mark_rect(opacity=0.08, color="gray")
            .encode(x=alt.X("x1:Q"), x2=alt.X2("x2:Q"))
        )

        _lines = (
            alt.Chart(_clip_df)
            .mark_line(strokeWidth=2)
            .encode(
                x=alt.X("rho:Q", title="Probability Ratio p"),
                y=alt.Y("value:Q", title="Surrogate Loss Contribution"),
                color=alt.Color(
                    "advantage:N",
                    scale=alt.Scale(range=["#e15759", "#4e79a7"]),
                    legend=alt.Legend(title="Advantage"),
                ),
                strokeDash=alt.StrokeDash(
                    "type:N",
                    scale=alt.Scale(domain=["Unclipped", "Clipped"], range=[[1, 0],
                    legend=alt.Legend(title="Line type"),
                ),
                tooltip=["rho:Q", "value:Q", "type:N", "advantage:N"],

```

```

    )
)

clipping_chart = (
    (_region_band + _lines)
    .properties(
        title="Surrogate Loss: Clipped vs Unclipped ( $\epsilon = 0.2$ )",
        width=580,
        height=320,
    )
    .interactive()
)

```

```
In [ ]: clipping_chart
```

The **shaded band** marks the $[1-\epsilon, 1+\epsilon] = [0.8, 1.2]$ region where clipping is inactive. Outside it the dashed line (clipped) flattens — removing the gradient signal that would push the policy too far from the old policy.

Algorithm Pseudocode

Algorithm: GRPO Fine-Tuning

Input : pre-trained LLM π_{ref} , dataset D , group size G ,
clip ϵ , KL weight β , learning rate η , steps T

Initialise $\pi_{\theta} \leftarrow \pi_{\text{ref}}$

```


for step = 1 ... T do
    Sample batch of prompts {q} from D
    for each prompt q do
        Sample G outputs {o_1,...,o_G} ~  $\pi_{\theta_{\text{old}}}(\cdot | q)$ 
        Compute rewards {r_1,...,r_G} (verifier / reward
model)
         $\bar{r} \leftarrow \text{mean}(\{r_i\})$ ,  $\sigma_r \leftarrow \text{std}(\{r_i\})$ 
         $A_i \leftarrow (r_i - \bar{r}) / (\sigma_r + \epsilon_{\text{stab}})$  # group
advantage
    end for
    for each (q, o_i, A_i) do
        for each token t in o_i do
             $\rho_{\{i,t\}} \leftarrow \pi_{\theta}(o_{\{i,t\}} | q, o_{\{i,<t\}}) /$ 
 $\pi_{\theta_{\text{old}}}(o_{\{i,t\}} | q, o_{\{i,<t\}})$ 
             $L_{\{i,t\}} \leftarrow \min(\rho_{\{i,t\}} \cdot A_i,$ 
clip( $\rho_{\{i,t\}}, 1-\epsilon, 1+\epsilon$ )  $\cdot A_i)$ 
        end for
    end for
    KL  $\leftarrow D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\text{ref}})$  # estimated over current batch
    J  $\leftarrow \text{mean}_i \text{mean}_t L_{\{i,t\}} - \beta \cdot \text{KL}$ 
     $\theta \leftarrow \theta + \eta \cdot \nabla_{\theta} J$ 
     $\theta_{\text{old}} \leftarrow \theta$  # update reference for next
step


```

end for

Return π_θ

GRPO Training Flowchart





Convergence Simulation: GRPO vs PPO

```

In [ ]: def _logistic(x: np.ndarray, L: float, k: float, x0: float) -> np.ndarray:
        return L / (1.0 + np.exp(-k * (x - x0)))

_steps = np.arange(1, 101)
_rng = np.random.default_rng(42)

_grpo_acc = _logistic(_steps, L=88.5, k=0.08, x0=40) + _rng.normal(0, 0.6, 100)
_ppo_acc = _logistic(_steps, L=83.0, k=0.06, x0=55) + _rng.normal(0, 0.6, 100)

_grpo_acc = np.clip(_grpo_acc, 0, 100)
_ppo_acc = np.clip(_ppo_acc, 0, 100)

_conv_df = pd.DataFrame({
    "Step": np.tile(_steps, 2),
    "Accuracy (%)": np.concatenate([_grpo_acc, _ppo_acc]),
    "Algorithm": ["GRPO"] * 100 + ["PPO"] * 100,
})
  
```


```
convergence_chart = (  
    alt.Chart(_conv_df)  
    .mark_line(point=False, strokeWidth=2)  
    .encode(  
        x=alt.X("Step:Q", title="Training Steps"),  
        y=alt.Y("Accuracy (%)ate="Accuracy (%)", scale=alt.Scale(domain=  
            color=alt.Color(  
                "Algorithm:N",  
                scale=alt.Scale(domain=["GRPO", "PPO"], range=["#f28e2b", "#4e79a2"  
            ),  
        tooltip=["Step:Q", "Accuracy (%)ate="Accuracy (%)", "Algorithm:N"],  
    )  
    .properties(  
        title="Simulated Convergence: GRPO vs PPO (synthetic data)",  
        width=580,  
        height=320,  
    )  
    .interactive()  
)
```

In []: convergence_chart

Algorithm Comparison Table

Algorithm	Advantage Estimate	Value Critic?	On-policy?	Notes
PPO	GAE (multi-step TD)	✔ Yes	Yes	Stable; high memory
TRPO	GAE	✔ Yes	Yes	Exact KL constraint; slow
DPO	Implicit (preference pairs)	✘ No	No	Offline; no reward model needed
GRPO	Group-normalised scalar	✘ No	Yes	Efficient; outcome-supervised
Hybrid-GRPO	Group + partial TD	Partial	Yes	Balances bias/variance
RPO	Group + robust penalty	✘ No	Yes	Worst-group optimisation
RE-PO	Group + entropy bonus	✘ No	Yes	Prevents entropy collapse
Group-Robust PO	Min-group advantage	✘ No	Yes	CVaR-style fairness objective

Timeline: Key Milestones



Hyperparameter Guidance

Hyperparameter	Typical Range	Effect
Group size G	4 - 16	Larger $G \rightarrow$ lower variance advantage; more compute per step
Clip ϵ	0.1 - 0.3	Smaller $\rightarrow$ more conservative updates; larger $\rightarrow$ faster but risky
KL weight β	0.001 - 0.1	Larger $\rightarrow$ stays closer to reference policy; smaller $\rightarrow$ more exploration
Learning rate	1e-6 - 5e-6	Use warmup; too large causes instability with clipping
Optimisation steps per batch	1 - 4	> 1 risks off-policy drift since data is from π_{old}
Reward normalisation	Per-group	Must normalise <i>within group</i> , not globally, to preserve relative signal

Practical Tips

- Start with $G = 8, \epsilon = 0.2, \beta = 0.01$ as a baseline.
- If reward variance is very low (all-correct or all-wrong groups) increase G .
- If the model diverges from the reference too quickly, increase β .
- Monitor token-level entropy; a sharp drop signals entropy collapse $\rightarrow$ use Dr. GRPO or RE-PO.
- For math/code tasks with a symbolic verifier the KL term may be reduced since the verifier already provides clean reward signal.

Empirical Results: DeepSeekMath Benchmarks

```
In [ ]: _results_df = pd.DataFrame({
    "Task": ["GSM8K", "GSM8K", "MATH", "MATH", "CMATH", "CMATH"],
    "Method": ["PPO", "GRPO", "PPO", "GRPO", "PPO", "GRPO"],
    "Accuracy (%)": [82.9, 88.2, 46.8, 51.7, 84.6, 88.8],
})

_bar = (
```

```

alt.Chart(_results_df)
  .mark_bar()
  .encode(
    x=alt.X("Task:N", title=None, axis=alt.Axis(labelFontSize=13)),
    xOffset=alt.XOffset("Method:N"),
    y=alt.Y(
      "Accuracy (%) :Q",
      title="Accuracy (%)",
      scale=alt.Scale(domain=[40, 95]),
    ),
    color=alt.Color(
      "Method:N",
      scale=alt.Scale(domain=["PP0", "GRPO"], range=["#4e79a7", "#f28e2b"]),
      legend=alt.Legend(title="Method"),
    ),
    tooltip=["Task:N", "Method:N", alt.Tooltip("Accuracy (%) :Q", format="%.1f")]
  )

_text = (
  alt.Chart(_results_df)
  .mark_text(dy=-6, fontSize=11, fontWeight="bold")
  .encode(
    x=alt.X("Task:N"),
    xOffset=alt.XOffset("Method:N"),
    y=alt.Y("Accuracy (%) :Q"),
    text=alt.Text("Accuracy (%) :Q", format="%.1f"),
    color=alt.Color("Method:N", scale=alt.Scale(domain=["PP0", "GRPO"], range=["#4e79a7", "#f28e2b"])),
  )
)

empirical_chart = (
  (_bar + _text)
  .properties(
    title="GRPO vs PP0 – DeepSeekMath Benchmarks",
    width=480,
    height=340,
  )
  .interactive()
)

```

In []: empirical_chart

Empirical Results Table

Results from DeepSeekMath (Shao et al., 2024) on the 7B model:

Benchmark	SFT Baseline	PP0	GRPO	GRPO Δ vs PP0
GSM8K	74.8%	82.9%	88.2%	+5.3 pp
MATH	35.9%	46.8%	51.7%	+4.9 pp
CMATH	76.6%	84.6%	88.8%	+4.2 pp

Benchmark	SFT Baseline	PPO	GRPO	GRPO Δ vs PPO
OCW Courses	12.8%	15.1%	16.6%	+1.5 pp

GRPO consistently outperforms PPO across all math benchmarks while using **half the GPU memory** (no value model) and achieving faster wall-clock training time.

Open Problems & Future Directions

Known Failure Modes

1. **Entropy Collapse** — Extended GRPO training can reduce output diversity, causing the model to repeat a narrow set of solution patterns. *Mitigation:* RE-PO entropy bonus; DAPO entropy-regularised objective.
2. **Reward Hacking** — With weak verifiers the model can exploit loopholes (e.g. formatting tricks) rather than learning genuine reasoning. *Mitigation:* Robust reward shaping; process reward models (PRMs).
3. **Off-Policy Drift** — Multiple gradient steps per sample accumulate importance-sampling error. *Mitigation:* Limit to 1-2 steps; re-sample frequently.
4. **Group Homogeneity** — If all G outputs receive identical rewards ($\sigma_r \approx 0$) the advantage is ill-defined and the update is zero. *Mitigation:* Adaptive G ; temperature annealing to encourage diversity.

Future Directions

- **Process-level GRPO:** Reward intermediate reasoning steps (PRMs) rather than only final answers to provide denser signal.
- **Multi-turn GRPO:** Extend to dialogue settings with session-level rewards.
- **Continual GRPO:** Online adaptation where the dataset itself grows from model interactions.
- **Theoretical analysis:** Tight convergence bounds for the critic-free group estimator vs GAE; characterising the bias-variance frontier as G varies.
- **Fairness-aware variants:** Group-Robust PO minimises worst-group error, relevant for demographic-fairness in alignment.

Further Reading

Resource	Type	Notes
Shao et al. (2024) — <i>DeepSeekMath</i>	Paper	Original GRPO paper; math RL fine-tuning

Resource	Type	Notes
DeepSeek-R1 Technical Report (2025)	Report	GRPO for long CoT reasoning at scale
Schulman et al. (2017) — <i>PPO</i>	Paper	Foundational clipped surrogate objective
Schulman et al. (2015) — <i>TRPO</i>	Paper	Trust-region policy optimisation
Rafailov et al. (2023) — <i>DPO</i>	Paper	Offline preference optimisation baseline
Ouyang et al. (2022) — <i>InstructGPT</i>	Paper	RLHF pipeline with PPO
Dr. GRPO (2024)	Paper	Addresses entropy collapse in GRPO
DAPO (2025)	Paper	Dynamic sampling + entropy regularisation
Group-Robust PO (2024)	Paper	CVaR-style fairness-aware variant
Hugging Face TRL docs	Docs	Open-source GRPO trainer implementation